

Ensemble Learning

Fundamentals of Machine Learning

Xufeng Zhang

Inria
xufeng.zhang@inria.fr

December 11, 2025

Outline

- 1 Motivation and Basics
- 2 Decision Trees Recap
- 3 Bagging and Random Forests
- 4 Boosting and Gradient Boosting
- 5 Modern Boosting Libraries
- 6 Stacking and Other Ensembles
- 7 Practical Guidance
- 8 Summary

What is Ensemble Learning?

- **Ensemble learning:** combine multiple base models to obtain a better predictive model.
- Key intuition:
 - Different models make different errors.
 - If their errors are not perfectly correlated, averaging can reduce overall error.
- Works for both classification and regression.
- Often one of the most effective ways to improve performance without changing features.

Voting Classifiers

Hard Voting

- Each classifier outputs a class label.
- Final prediction is the **majority class**.
- Ties are broken by a fixed class order or additional rules.

Soft Voting

- Each classifier outputs class probabilities $p_{m,c}$.
- Weighted average:

$$\hat{y} = \arg \max_c \sum_{m=1}^M w_m p_{m,c},$$

where $\sum_m w_m = 1$.

- More informative than hard voting if probabilities are well calibrated.

Why Does Ensembling Work?

- Reduces **variance**: unstable models (e.g., deep trees) change a lot with data perturbations.
- Can reduce **bias**: sequential methods (Boosting) focus on residual errors.
- Main requirements:
 - Base models must be **better than random**.
 - Base models must be **diverse** (uncorrelated errors).
- Diversity can come from:
 - Different training samples.
 - Different model architectures.
 - Different hyperparameters or feature subsets.

Bias–Variance Perspective

- Prediction error (expected squared error):

$$\mathbb{E}[(Y - \hat{f}(X))^2] = \underbrace{\text{Bias}^2}_{\text{systematic error}} + \underbrace{\text{Variance}}_{\text{sensitivity to data}} + \text{Noise}.$$

- **Bagging / Random Forest:**

- Typically use high-variance, low-bias models.
- Aim: reduce variance by averaging.

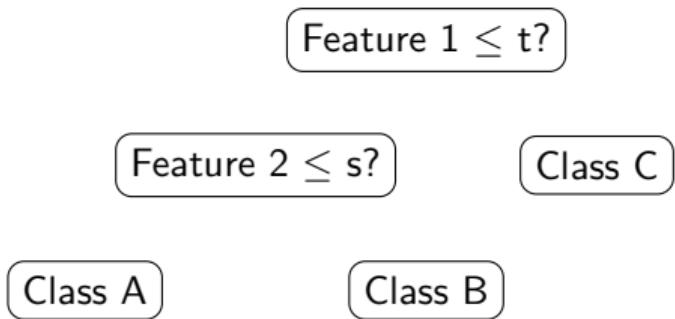
- **Boosting:**

- Typically use low-variance, high-bias base learners.
- Aim: reduce bias by combining many weak learners.

Decision Trees: Representation

- Tree structure with **internal nodes** = tests on features.
- **Leaves** correspond to regions of feature space.
- Classification: leaf stores class probabilities (relative frequencies).
- Regression: leaf stores the mean target value.
- Advantage: easy to visualize and interpret.

Example: Simple Classification Tree



- Each root-to-leaf path defines a **decision rule**.
- Trees partition the space using axis-aligned splits.

Splitting Criteria and Impurity

Consider a leaf with class probabilities \hat{p}_c .

Gini Impurity

$$G(\hat{\mathbf{p}}) = \sum_c \hat{p}_c(1 - \hat{p}_c) = 1 - \sum_c \hat{p}_c^2.$$

Entropy

$$H(\hat{\mathbf{p}}) = - \sum_c \hat{p}_c \log_2 \hat{p}_c.$$

- Splits are chosen to maximize **information gain**.
- Information gain = impurity(parent) – weighted impurity(children).

Regression Trees

- For regression, each leaf predicts:

$$\mu_L = \frac{1}{|L|} \sum_{i \in L} y_i.$$

- Split chosen to minimize squared error in leaves:

$$\sum_L \sum_{i \in L} (y_i - \mu_L)^2.$$

- Predictions are piecewise constant: non-smooth and cannot extrapolate.
- Still powerful as base learners inside ensembles.

Feature Importance from Trees

- Trees can compute **impurity-based feature importance**.
- Importance of feature j :
 - Sum of impurity reductions over all splits on feature j ,
 - Weighted by the number of samples at each node.
- In Random Forests / Gradient Boosting:
 - Importances are averaged over many trees.
 - More robust than using a single tree.

Under- and Overfitting with Trees

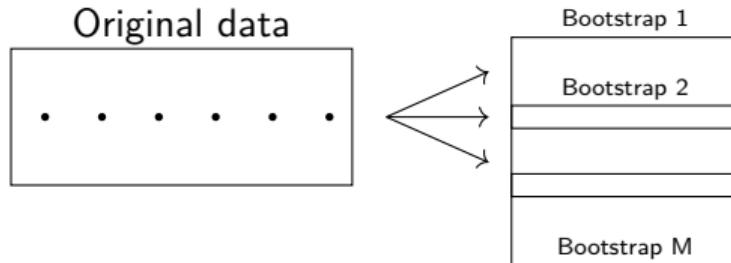
- **Shallow trees:**
 - High bias, low variance.
 - Capture only simple patterns (underfitting).
- **Deep trees:**
 - Low bias, high variance.
 - Fit training data very well (risk of overfitting).
- Trees are ideal base learners for ensembles:
 - Deep trees for Bagging / Random Forest (variance reduction).
 - Shallow trees (stumps) for Boosting (bias reduction).

Bagging: Bootstrap Aggregating

- Given training set D of size n .
- For $m = 1, \dots, M$:
 - Draw bootstrap sample D_m of size n with replacement.
 - Train base model h_m on D_m .
- For a new point x :
 - Classification: majority vote or soft vote.
 - Regression: average prediction $\hat{y}(x) = \frac{1}{M} \sum_m h_m(x)$.
- Reduces variance by averaging many high-variance estimators.

Bootstrap Sampling Illustration

- Each bootstrap sample leaves out about $\approx 36\%$ of instances (out-of-bag).
- Different samples lead to different fitted trees.
- Averaging the predictions smooths the decision boundary.



Random Forests

- Extension of Bagging using **randomized trees**.
- At each split:
 - Consider only a random subset of features of size `max_features`.
 - Choose the best split among those features.
- Further de-correlates trees:
 - Trees become more diverse.
 - Ensemble variance decreases more effectively.
- Extremely randomized trees (ExtraTrees):
 - Also randomize split thresholds.
 - Faster and sometimes more regularized.

Effect on Bias and Variance

- Increasing the number of trees M :
 - **Variance** decreases and stabilizes.
 - **Bias** is mostly unchanged (trees are still high-capacity).
- With very many trees:
 - Decision boundary becomes smoother.
 - Too much smoothing can slightly increase bias.
- In practice:
 - Use large M (e.g., 100–1000) until performance plateaus.
 - Training/prediction time is the main limitation.

Important Hyperparameters (Random Forest)

- `n_estimators`:
 - Number of trees; higher is usually better.
 - Diminishing returns after some point.
- `max_features`:
 - Typical defaults: \sqrt{p} for classification, $\log_2(p)$ or p for regression.
 - Smaller values increase decorrelation but may increase bias.
- Tree parameters:
 - `max_depth`, `min_samples_split`, `min_samples_leaf`, etc.
 - Pre-pruning reduces tree size and training time.
- `bootstrap`, `oob_score` (out-of-bag error).

Out-of-Bag (OOB) Error

- For each tree, about 1/3 of training samples are **not** included in its bootstrap sample.
- These are **out-of-bag** (OOB) for that tree.
- For each training point:
 - Aggregate predictions over trees for which the point is OOB.
- Compute OOB error using these aggregated predictions.
- OOB error:
 - Comparable to cross-validation performance.
 - Saves computation during model selection.

Random Forests: Strengths and Weaknesses

Strengths

- Strong baseline; often high accuracy with minimal tuning.
- Handles heterogeneous feature types and scales.
- Provides feature importance estimates.
- Parallelizable across trees.

Weaknesses

- Less interpretable than single trees.
- Larger memory footprint, slower predictions.
- Not ideal for very high-dimensional sparse data (e.g., text).
- Probability estimates can be poorly calibrated.

Boosting: High-Level Idea

- Build an ensemble sequentially.
- Each new model focuses on examples that previous models got wrong.
- Final model is an additive combination:

$$F_M(x) = \sum_{m=1}^M \alpha_m h_m(x).$$

- Two major families:
 - **AdaBoost**: reweight samples based on previous errors.
 - **Gradient Boosting**: fit to residuals (gradient steps).

AdaBoost: Intuition

- Use simple base learners: usually decision stumps (depth-1 trees).
- Start with equal weights on all training examples.
- After each iteration:
 - Increase weights of misclassified examples.
 - Decrease weights of correctly classified examples.
- New weak learner is trained on this reweighted dataset.
- Hard examples get more influence over time.

AdaBoost: Algorithm (Binary Classification)

- ① Initialize sample weights $w_i^{(1)} = 1/N$.
- ② For $m = 1, \dots, M$:
 - ① Train weak learner h_m using weights $w^{(m)}$.
 - ② Compute weighted error:

$$\varepsilon_m = \frac{\sum_i w_i^{(m)} \mathbf{1}\{h_m(x_i) \neq y_i\}}{\sum_i w_i^{(m)}}.$$

- ③ Compute learner weight:

$$\alpha_m = \frac{1}{2} \log \frac{1 - \varepsilon_m}{\varepsilon_m}.$$

- ④ Update sample weights:

$$w_i^{(m+1)} = w_i^{(m)} \exp(-\alpha_m y_i h_m(x_i)),$$

then renormalize.

- ⑤ Final prediction: $\text{sign}(\sum_m \alpha_m h_m(x))$.

AdaBoost: Bias–Variance Behaviour

- Weak learners have high bias, low variance.
- AdaBoost reduces bias by combining many such learners.
- Initially:
 - Strong improvements in both training and test error.
- If we keep boosting:
 - Training error can reach zero.
 - Variance may increase; risk of overfitting.
- In practice:
 - Control complexity via `n_estimators` and `learning_rate`.

Gradient Boosting: General Framework

- Want to minimize a differentiable loss $L(y, F(x))$.
- Initialize model $F_0(x)$ (e.g., constant).
- For $m = 1, \dots, M$:
 - ① Compute pseudo-residuals:

$$r_i^{(m)} = - \left. \frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right|_{F=F_{m-1}}.$$

- ② Fit base learner $h_m(x)$ to predict $r^{(m)}$.
- ③ Choose step size η (learning rate).
- ④ Update:

$$F_m(x) = F_{m-1}(x) + \eta h_m(x).$$

- Base learners: typically shallow regression trees.

Gradient Boosting for Regression

- Loss: squared error $L(y, F(x)) = \frac{1}{2}(y - F(x))^2$.
- Pseudo-residuals:

$$r_i^{(m)} = y_i - F_{m-1}(x_i),$$

i.e., the regression residuals.

- Each tree approximates the residuals of the current ensemble.
- Final prediction:

$$F_M(x) = F_0(x) + \eta \sum_{m=1}^M h_m(x).$$

- Early stopping: monitor validation error and stop when it stops improving.

Gradient Boosting for Classification

- Often implemented via **logistic loss**.
- For binary labels $y_i \in \{0, 1\}$, model log-odds:

$$F(x) \approx \log \frac{p(y = 1 \mid x)}{p(y = 0 \mid x)}.$$

- Pseudo-residuals:

$$r_i^{(m)} = y_i - p_{m-1}(x_i),$$

where $p_{m-1} = \sigma(F_{m-1}(x_i))$ is the predicted probability.

- Trees fit these residuals; ensemble updates log-odds.
- Final probability via sigmoid:

$$p_M(x) = \sigma(F_M(x)).$$

Gradient Boosting: Hyperparameters

- `n_estimators`:
 - Number of boosting iterations.
 - Too many iterations can overfit.
- `learning_rate (η)`:
 - Smaller η requires more trees but can improve generalization.
 - Common range: 0.01 to 0.3.
- Tree complexity:
 - `max_depth` or `max_leaf_nodes`.
 - Usually 2–5 for base learners.
- Regularization:
 - Subsample rows or columns.
 - Early stopping via `n_iter_no_change`.

Gradient Boosting: Pros and Cons

Pros

- Among the most accurate off-the-shelf models.
- Handles different feature types and scales.
- Flexible via choice of loss function.

Cons

- Training is sequential and harder to parallelize.
- Sensitive to hyperparameters; more tuning required than Random Forests.
- Can overfit if `n_estimators` is large and regularization is weak.

Extreme Gradient Boosting (XGBoost)

- Highly optimized gradient boosting implementation.
- Key features:
 - Uses second-order derivatives (Hessian) for faster convergence.
 - Regularization on tree weights and leaf scores.
 - Efficient handling of sparse inputs and missing values.
 - Histogram-based split finding (feature binning).
- Supports parallel and distributed training.
- Widely used in machine learning competitions and industry.

LightGBM and CatBoost

LightGBM

- Gradient boosting framework from Microsoft.
- Uses gradient-based one-side sampling:
 - Keep all instances with large gradients.
 - Randomly sample among those with small gradients.
- Leaf-wise tree growth with depth constraints.
- Efficient handling of large-scale datasets.

CatBoost

- Gradient boosting framework optimized for categorical features.
- Uses ordered target encoding and symmetric trees.
- Strong default settings; less need for heavy tuning.
- Supports monotonicity constraints for selected features.

Stacking: Idea

- Combine **heterogeneous** base models (e.g., RF, GBM, SVM, NN).
- For each base model h_j , obtain predictions $\hat{y}_j(x)$.
- Train a **meta-model** (stacker) on these predictions:

$$z(x) = [\hat{y}_1(x), \dots, \hat{y}_K(x)], \quad \tilde{y} = g(z(x)).$$

- Often use cross-validation to create out-of-fold predictions for training the stacker.
- Popular stackers: linear models, logistic regression, gradient boosting.

Stacking: Design Considerations

- Base models should be:
 - **Strong** performers individually.
 - **Diverse** (different biases and errors).
- Avoid information leakage:
 - Use cross-validated predictions for meta-training.
 - Do not train meta-model on predictions from the same folds used to fit base models.
- Complexity and speed:
 - Stacking can be expensive at prediction time (many models).
 - Useful when accuracy is the main objective.

Other Ensemble Techniques

- **Hyper-ensembles:**
 - Same algorithm with different hyperparameter settings.
 - Example: multiple Random Forests with different depths.
- **Deep ensembles:**
 - Train multiple neural networks with different initializations or subsets.
 - Improve robustness and uncertainty estimates.
- **Bayesian model averaging:**
 - Average over models weighted by posterior probabilities.
 - More theoretical; often approximated by ensembles.
- **Cross-validation selection:**
 - Select the single best model via internal cross-validation (not strictly an ensemble).

When to Use Which Ensemble?

- **Random Forests:**
 - Strong, robust baseline.
 - Limited need for tuning; good with small to medium tabular data.
- **Gradient Boosting / XGBoost / LightGBM / CatBoost:**
 - Often best performance on structured/tabular data.
 - Requires more tuning; good when accuracy is critical.
- **AdaBoost:**
 - Useful for simple weak learners.
 - Historically important; less used than modern GBMs.
- **Stacking:**
 - Final refinement when diverse high-performing models are available.

Calibration, Warm Starts, and Early Stopping

- **Calibration:**

- Tree ensembles often produce poorly calibrated probabilities.
- Use methods like isotonic regression or Platt scaling.

- **Warm starting:**

- Continue training by adding trees to an existing ensemble.
- Useful for incremental improvement on similar data.

- **Early stopping:**

- Monitor validation loss; stop adding trees when performance stops improving.
- Avoids overfitting and saves computation.

Typical Ensemble Workflow

- ① Start with a simple model (e.g., logistic regression, decision tree) for baseline.
- ② Train a Random Forest; tune `n_estimators`, `max_depth`, `max_features`.
- ③ Train a Gradient Boosting or XGBoost model with early stopping.
- ④ Compare performance on validation / cross-validation.
- ⑤ Optionally:
 - Calibrate probabilities.
 - Build a small stacking ensemble from the best models.
- ⑥ Evaluate final ensemble on a held-out test set.

Algorithm Overview

Method	Base Learners	Loss	Optimization
Decision Trees	Single tree	Gini / Entropy / MSE	Greedy splitting
Random Forest	Many trees	Gini / Entropy / MSE	Bagging + feature subsampling
AdaBoost	Stumps / shallow trees	Exponential	Greedy reweighting
Grad. Boosting (Reg.)	Regression trees	MSE	Gradient descent on residuals
Grad. Boosting (Clf.)	Regression trees	Log loss	Gradient descent on log-odds
XGBoost / LightGBM / CatBoost	Small trees	Task-specific	2nd order gradients + regularization
Stacking	Any models	Task-specific	Meta-model on predictions

Key Takeaways

- Ensembles leverage multiple models to reduce error via bias and variance control.
- Bagging / Random Forest:
 - Variance reduction by averaging many overfitting models.
- Boosting (AdaBoost, Gradient Boosting):
 - Bias reduction by focusing on hard examples or residuals.
- Modern boosting libraries (XGBoost, LightGBM, CatBoost) are state-of-the-art for tabular data.
- Stacking combines heterogeneous models using a learned combiner.
- Always consider computational cost, interpretability, and deployment constraints.

References

- Lecture 4: *Ensemble Learning — ML Engineering*, Joaquin Vanschoren.
- L. Breiman, “Random Forests,” Machine Learning, 2001.
- J. Friedman, T. Hastie, R. Tibshirani, “The Elements of Statistical Learning.”
- T. Chen and C. Guestrin, “XGBoost: A Scalable Tree Boosting System,” KDD 2016.